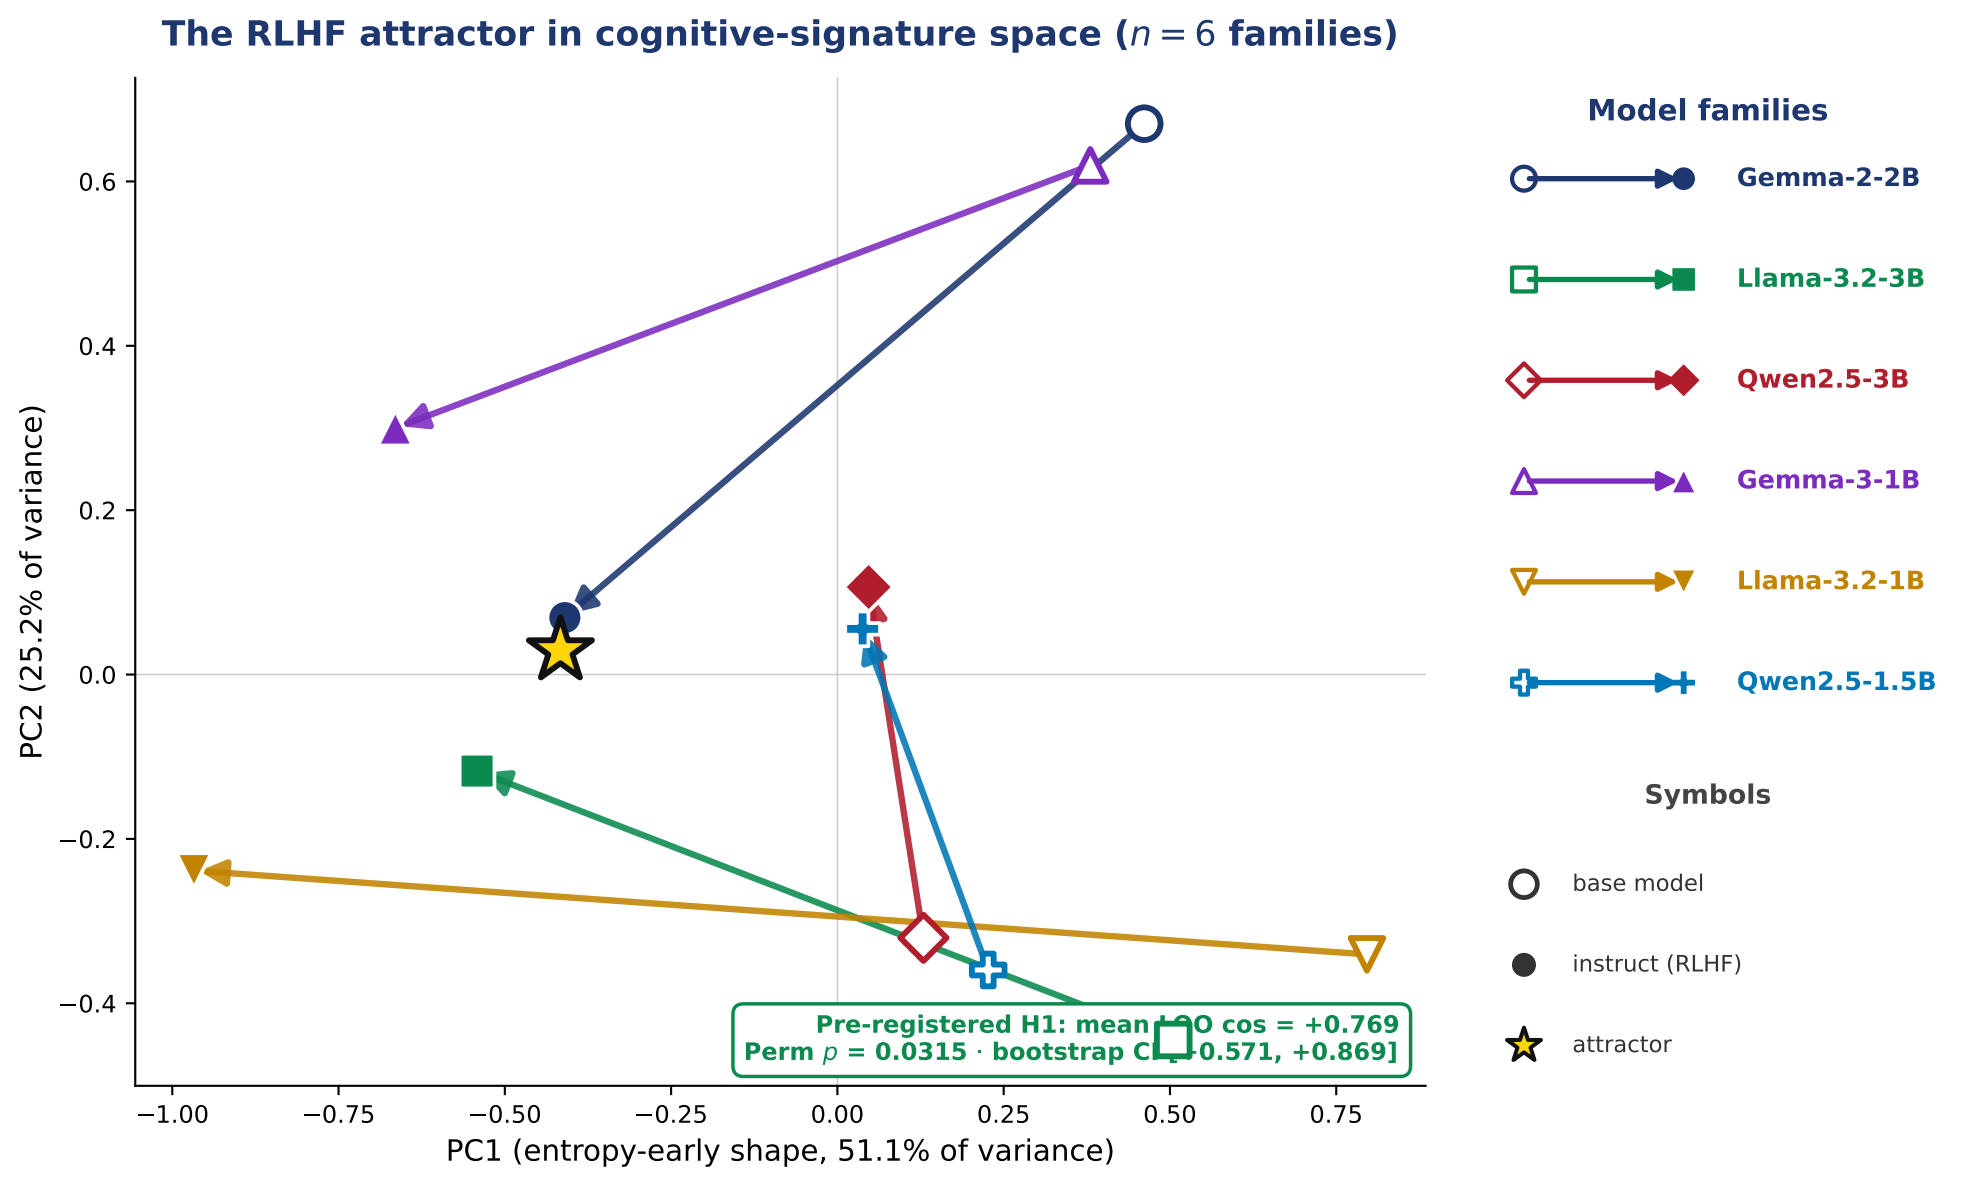
The RLHF attractor in cognitive-signature space ($n = 6$ families)
PC2 (25.2% of variance)
PC1 (entropy-early shape, 51.1% of variance)
0.6
0.4
0.2
0.0
−0.2
−0.4
−1.00
−0.75
−0.50
−0.25
0.00
0.25
0.50
0.75
Pre-registered H1: mean LOO cos = +0.769
Perm p = 0.0315 · bootstrap CI [+0.571, +0.869]
Model families
Gemma-2-2B
Llama-3.2-3B
Qwen2.5-3B
Gemma-3-1B
Llama-3.2-1B
Qwen2.5-1.5B
Symbols
base model
instruct (RLHF)
attractor